

Introduction to Monte Carlo Simulations

Scenarios, Breakthroughs, and Applications

Minwoo Yoo

Department of Economics
George Washington University

Econ 8376 - Spring 2026

The Goal: Knowing Your Estimator

We build an estimator $\hat{\theta}$ to learn about a parameter θ .

Before we trust $\hat{\theta}$, we need to know its properties:

- Is it **unbiased**? ($E[\hat{\theta}] = \theta$?)
- What is its **variance**?
- What is its **distribution**? (Normal? Chi-squared?)

Usually, we use **Asymptotic Theory** (Math) to answer these.

But what if the math fails?

Scenario A: The "Small Sample" Anxiety

The Situation:

- You derived that $\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} N(0, V)$ as $n \rightarrow \infty$.
- But in your actual research, you only have **30 observations** ($n = 30$).

The Obstacle:

- Is $n = 30$ large enough? Is the approximation accurate?
- Or is the finite-sample distribution actually skewed?

Scenario A: The "Small Sample" Anxiety

The Situation:

- You derived that $\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} N(0, V)$ as $n \rightarrow \infty$.
- But in your actual research, you only have **30 observations** ($n = 30$).

The Obstacle:

- Is $n = 30$ large enough? Is the approximation accurate?
- Or is the finite-sample distribution actually skewed?

The Breakthrough

Don't guess. Simulate.

- ① Create 10,000 artificial worlds with $n = 30$.
- ② Calculate $\hat{\theta}$ in each world.
- ③ **Plot the histogram.** If it's not Normal, the asymptotic theory is misleading you.

Scenario B: The "Impossible Integral"

The Situation (Structural Models / Bayesian):

- You need to calculate an expected value (e.g., Market Shares in BLP).
- This involves a high-dimensional integral:

$$\sigma_j(p) = \int \cdots \int P_j(p, \nu) f(\nu) d\nu$$

The Obstacle:

- There is **no closed-form solution**. You cannot solve it by hand.

Scenario B: The "Impossible Integral"

The Situation (Structural Models / Bayesian):

- You need to calculate an expected value (e.g., Market Shares in BLP).
- This involves a high-dimensional integral:

$$\sigma_j(p) = \int \cdots \int P_j(p, \nu) f(\nu) d\nu$$

The Obstacle:

- There is **no closed-form solution**. You cannot solve it by hand.

The Breakthrough (Monte Carlo Integration)

Transform Calculus into Arithmetic. By the Law of Large Numbers (LLN), sample average converges to expectation:

$$\int g(\nu) f(\nu) d\nu \approx \frac{1}{R} \sum_{r=1}^R g(\nu_r)$$

Just draw random ν_r and take the average.

Scenario C: The "Black Box" Estimator

The Situation:

- Your estimator $\hat{\theta}$ is the result of a complex algorithm or optimization (e.g., a complicated GMM or ML estimator).
- You cannot easily derive the formula for its Variance.

The Obstacle:

- How do I report the Standard Error?

Scenario C: The "Black Box" Estimator

The Situation:

- Your estimator $\hat{\theta}$ is the result of a complex algorithm or optimization (e.g., a complicated GMM or ML estimator).
- You cannot easily derive the formula for its Variance.

The Obstacle:

- How do I report the Standard Error?

The Breakthrough

Empirical Variance.

- Run the algorithm 1,000 times on simulated data.
- Calculate the variance of the 1,000 results.
- That is your approximation of the true variance.

So, What is This?

This approach of "solving problems by generating random numbers" is called:

Monte Carlo Simulation

It replaces **analytical derivation** (which might be hard or impossible) with **computational repetition** (which is easy for computers).

The Experiment Design: Testing the CLT

The Question: Standard theory (CLT) says \bar{Y} should be Approx. Normal even if N is finite. Is $N = 25$ "large enough" for this to hold?

The Blueprint: To verify this, we will act as the **Data Generating Process (DGP)**.

- ① **Assumption:** The true population is Standard Normal ($Y \sim N(0, 1)$).
- ② **Action:** Draw $N = 25$ random observations from this population.
- ③ **Calculation:** Compute the sample mean \bar{Y} .
- ④ **Repetition:** Do this 100,000 times to approximate the distribution of \bar{Y} .
- ⑤ **Verification:** Check if the resulting histogram looks like a Bell Curve.

Example: Setup the Estimator

Let's translate our blueprint into Python.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # 1. Define the "Truth" (DGP Parameters)
5 n = 25          # Small sample size
6 muY = 0         # True Mean
7 sigY = 1         # True Std Dev
8
9 # 2. Define the Estimator (The "Rule")
10 def get_sample_mean(n, muY, sigY):
11     # Generate data from the true distribution
12     Y = sigY * np.random.randn(n) + muY
13
14     # Return the estimate (Sample Mean)
15     return Y.mean()
```

Example: The Simulation Loop

```
1 # 3. The Monte Carlo Loop (Repeat 100,000 times)
2 MC = 100000
3 sample_means = np.zeros(MC) # Container for results
4
5 for i in range(MC):
6     # Store the result of each experiment
7     sample_means[i] = get_sample_mean(n, muY, sigY)
8
9 # 4. Analysis
10 print("Exp. Value:", sample_means.mean()) # Check Unbiasedness
11 print("Variance:", sample_means.var())    # Check Efficiency
```


Example: Visualization Code

Finally, we draw the histogram to see the distribution.

```
1 # 5. Plot the sampling distribution
2 plt.hist(sample_means, bins=50, color='skyblue', edgecolor='black')
3
4 plt.title("Sampling Distribution of Sample Mean (N=25)")
5 plt.xlabel("Sample Mean Value")
6 plt.ylabel("Frequency")
7
8 # Save or Show the plot
9 plt.show()
```

The Result: Visual Proof

- We generated 100,000 **Sample Means**.
- The histogram shows the **Sampling Distribution**.
- **Conclusion:** Even with $N = 25$, the distribution of \bar{Y} looks fairly Normal (Bell-shaped).
- The CLT works reasonably well here!



Comparison: Monte Carlo vs. Bootstrap

	Monte Carlo	Bootstrap
When to use?	To study properties of an estimator (e.g., Does \bar{Y} work well?).	To calculate SEs/CIs for <i>your specific dataset</i> .
The "World"	We create the world (DGP is known).	The sample is the world (DGP is unknown).

Summary & Best Practices:

- Use MC when math is too hard or N is too small.
- Always set a seed ('np.random.seed') for reproducibility.
- Ensure enough repetitions ($R \geq 1000$).